

Real-time Domain Adaptation in Semantic Segmentation*

*Note: Sub-titles are not captured in Xplore and should not be used

1st Berardo Nicholas
dept. Computer Engineering
Politecnico di Torino
Turin, Italy
s319349@studenti.polito.it

2nd Cardona Riccardo
dept. Computer Engineering
Politecnico di Torino
Turin, Italy
s319441@studenti.polito.it

3rd De Marco Alessandro
dept. Computer Engineering
Politecnico di Torino
Turin, Italy
sXXXXXXX@studenti.polito.it

Abstract—We use an efficient structure named Short-Term Dense Concatenate network (STDC network) for the semantic segmentation task. This structure reduce the dimension of feature maps and use the aggregation of them for image representation, then use a Detail aggregation module for producing the low-level features. Finally these two are merged to produce the segmentation result. We test this model on Cityscapes and GTA V, following the evaluation of the domain shift between GTA V and Cityscapes and finally we implement an unsupervised adversarial domain adaptation method used for reducing the domain shift. We also show the result for the STDC network in term of mIoU and the result for the domain adaptation.

I. INTRODUCTION

Semantic Segmentation is a topic in computer vision that aims at assigning a label to each pixel of the image. This is used in many fields such as autonomous vehicle, video surveillance and robot sensing. There are a lot of models that can achieve good accuracy. For real-time semantic segmentation some models choose lightweight backbones for having an increase of performance but with an acceptable loss in accuracy. For this reason some new methods were investigated, like feature fusion or aggregation modules. Other models reduce the input image size but this can result in a bad accuracy around boundaries and small object.

STDC net [1] uses the first approach. Fig. 1 shows how the image is encoded in different scales. The kernel size is also reduced to speed-up the performance but with an acceptable loss in accuracy. Then a Detail Guidance is used to learn the space details instead of using a Spatial Path as in BiSeNet [2].

The next step is domain adaptation. A model trained on a certain dataset may not generalize on an unseen dataset ending in poor performance. This is caused by the domain shift between the source (training) and target (test) dataset, for example different cities, weather and lighting conditions. Domain adaptation methods are used to close the gap between source and target domains. In this paper we use an adversarial domain adaptation method [3] composed of a segmentation model to predict the output and a discriminator to predict if the input is from the source or target domain. The goal is to generate a segmentation output from the segmentation

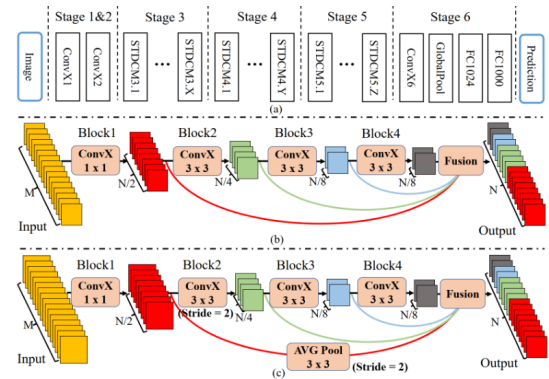


Fig. 1. STDC network.

part that fool the discriminator, meaning that the segmentation output is similar between source and target domains. We show experiments done on the adaptation between GTAV and Cityscapes.

Our contributions can be summarized as follows: First we build the STDC network and train it on Cityscapes and test it again on Cityscapes. Second we apply the same idea over the GTAV dataset, so train on GTAV and test on GTAV. Third we compute the domain shift between GTAV (source) and Cityscapes (target) domains firstly in vanilla, then with some augmentation on GTAV. Fourth we implement the adversarial domain adaptation method and test it with GTAV as source domain and Cityscapes as target domain. Lastly we apply some real-time semantic segmentation method to our discriminator function to make it faster.

Here a look to our implementation:
<https://github.com/Riden15/Real-time-Domain-Adaptation-in-Semantic-Segmentation>

II. RELATED WORK

TODO Related work

III. METHODS

We first introduce the Short-Term Dense Concatenate network (STDC network) and how we used it with BiSeNet [2], then the unsupervised adversarial domain adaptation method.

STDC network [1] is represented in Fig. 1 (a). Stage 3,4 and 5 have a number of Short-Term Dense Concatenate Module (STDCM) where each module is composed of ConvX blocks Fig. 1 (b)(c). Each $ConvX_i$ is a block composed of one convolutional layer, one batch normalization layer and one ReLU activation layer. The ConvX layers filter the input into $N/2$, where N is the channel number of the STDC module. At the end we concatenate the output of each ConvX block as follow:

$$x_{output} = F(x_1, x_2, \dots, x_n)$$

where x_{output} is the STDC module output, F is the fusion operation, that in our case is the concatenation and x_1, x_2, \dots, x_n are the output of each $ConvX_i$ block.

This STDC network is then used as backbone for the context path of BiSeNet, in particular we use stage 3,4 and 5 to reduce the feature map to obtain large receptive field. Then a global average pooling is added on the tail of the context path. We also use Attention Refine Module (ARM) on Stage 4 and 5. Finally a Future Fusion Model (FFM) is used to fuse the 1/8 feature from Stage 3 with this last part.

The main idea of the unsupervised adversarial domain adaptation method [3] is to create a segmentation network that classify the source image (that as a label) and the target image (without annotations). Then those two predictions are given to the discriminator that has to distinguish whether the input is from source or target domain. To do so we need a loss from the discriminator to the segmentation network which encourage the segmentation to generate similar predictions. We use as segmentation network the BiSeNet developed above. The loss is

$$L(I_s, I_t) = L_{seg}(I_s) + \lambda_{adv} * L_{adv}(I_t)$$

where L_{seg} is the cross entropy loss for the prediction of the source domain, λ_{adv} is the weight used to balance the L_{adv} that is a BCEWithLogitLoss. We first forward I_s to the segmentation network to get the prediction P_s and the loss L_{seg} using the label of I_s . Then forward I_t to the segmentation network to get the prediction P_t and feed P_s to the discriminator and get the L_{adv} and optimize the segmentation network. Now let's optimize the discriminator computing first the loss on the GTA5 image with label 1 and finally computing the loss on the Cityscape image, with label 0.

Finally we also tried to apply some real-time semantic segmentation method to our discriminator to make it lighter and faster. We build a discriminator using depthwise separable convolutions, following MobileNets [6]. These are composed of two parts: depthwise convolutions and pointwise convolutions. Depthwise convolutions are used to filter each input channel, but they can only filter so we use pointwise convolutions (1 x 1 kernel size) to generate new features. Now our discriminator is based on blocks of depthwise convolution layer, leaky ReLU, pointwise convolution and again leaky ReLU.




Fig. 2. Loss step during training. On the X we have the epochs, on Y we have the loss. (1) is the gray curve, (2) is the light blue curve, (3) is the light blue curve again, (4) is the pink curve, (5) is the orange curve and (6) is the purple curve.




Fig. 3. mIoU during training. On the X we have the epochs, on Y we have the mIoU. (1) is the gray curve, (2) is the light blue curve, (4) is the pink curve, (5) is the orange curve and (6) is the purple curve.

IV. EXPERIMENTS

A. GTA5

GTA5 dataset [4] is composed of around 25K images with resolution 1914 x 1052 synthesized from the videogame based on Los Angeles. The dataset is also composed of the ground truth annotation for semantic segmentation. It is based on 19 categories, that are the same as in Cityscapes, so they are compatible. We use this dataset firstly for evaluating our model alone, so train and test on this dataset (obviously the test data is not seen during training). Then this dataset is used as source domain in domain adaptation.

B. Cityscapes

Cityscapes dataset [5] is composed of 5000 fine annotated images, split into training, validation and test sets with 2975, 500 and 1525 images respectively. The total number of classes is 30, but only 19 are used for semantic segmentation. The resolution of the images is 2048 x 1024, this is very high resolution, in fact we reduced it to 1024 x 512. This dataset is first used for evaluating our model alone, so train and test both on this dataset, then it is used as target domain for domain adaptation.

C. Experiments results

We have done the following experiments: (1) train on Cityscapes and test on Cityscapes, (2) train on GTA5 and test on GTA5, (3) train on GTA5 and test on Cityscapes without domain adaptation, (4) train on GTA5 and test on Cityscapes without domain adaptation with some augmentation on GTA5, (5) train on GTA5 and test on Cityscapes with domain adaptation and (6) train on GTA5 and test on Cityscapes with domain adaptation with a lighter discriminator. We use as optimizer the SGD (Stochastic Gradient Descent) with momentum set at 0.9 and weight decay at $5e^{-4}$. We use a batch size of 2 and poly learning rate where the learning rate is updated




Fig. 4. Precision during training. On the X we have the epochs, on Y we have the precision. (1) is the gray curve, (2) is the light blue curve, (4) is the pink curve, (5) is the orange curve and (6) is the purple curve.

TABLE I
TRAIN ON CITYSCAPES TRAIN SET AND TEST IT CITYSCAPES VAL SET

Cityscapes -> Cityscapes		
Accuracy (%)	mIoU (%)	Train time (avg per-epoch)
81.1	58.6	2.3 minutes

as follow $lr = init_lr * (1 - \frac{iter}{max_iter})^{power}$, where $init_lr$ is 0.01, $iter$ is the actual epoch, max_iter is the number of epochs and $power$ is 0.9. Data augmentation, when used, contains Color Jitter, Random Horizontal Flip, Random Crop and Normalization. The discriminator for the domain apadtation part is trained using Adam with a poly learning rate with the same parameter as before. We perform our experiments with those versions: **TODO!!!!!!!**

In Fig. 2 we can see the loss step during train time. The one that can achive a better low (lower) are the experiments where there isn't a domain shift. Same thing happend with the mean Intersection over Union (mIoU) in Fig. 3 and with the precision in Fig. 4.

(1) Here we train our model on Cityscapes train set and test it on Cityscapes validation set. Here there's no domain adaptation, so we expect to have high results. Table I shows the results.

(2) We train our model on GTA5 train set and test it on GTA5 validation set. Again we expect high results. Table II shows the results.

(3) Here we evaluate the domain shift between GTA5 and Cityscapes. We train on GTA5 train set and test on Cityscapes validation set. Here we don't have any domain adaptation method, so we expect bad results. Table III shows the results. We can see that the mIoU is low, this because the source domain and target domain are different.

(4) Again we evaluate the domain shift between GTA5 and Cityscapes, but this time with some augmentation on GTA5. Table IV shows the results and again the mIoU is low because of the domain shift, but higher than step (3). This thanks to the augmentation done on GTA5.

(5) Here we have domain adaptation [3]. So we train on GTA5 and test on Cityscapes, but this time we train also the discriminator with the segmentation network. Table V shows the results. We can see a slight improvement of performance in the mIoU of 3% with respect to (4).

(6) Finally we implement depthwise separable convolutions in the discriminator. Again same thing as (5) but with a lighter discriminator. Table VI shows the results.

TABLE II
TRAIN ON GTA5 TRAIN SET AND TEST IT GTA5 VAL SET

GTA5 -> GTA5		
Accuracy (%)	mIoU (%)	Train time (avg per-epoch)
80.8	62.3	3.2 minutes

TABLE III
TRAIN ON GTA5 TRAIN SET AND TEST IT CITYSCAPES VAL SET WITHOUT ADAPTATION

GTA5 -> Cityscapes		
Accuracy (%)	mIoU (%)	Train time (avg per-epoch)
60.1	24.6	3.2 minutes

V. CONCLUSION

In this paper we build a model based on BiSeNet [2] with a Short-Term Dense Concatenate network (STDC network) [1] as Context Path, based on different STDC modules. Then we implement an adversarial domain adaptation method [3] to reduce the domain shift between GTA5 and Cityscapes. From the experiments we can see that having a domain adaptation method can help in increasing the accuracy and mIoU. Finally we develop a lighter and faster method for discriminate the source and target data in the domain adaptation method, based on MobileNets [6].

ACKNOWLEDGMENT

The preferred spelling of the word “acknowledgment” in America is without an “e” after the “g”. Avoid the stilted expression “one of us (R. B. G.) thanks ...”. Instead, try “R. B. G. thanks...”. Put sponsor acknowledgments in the unnumbered footnote on the first page.

REFERENCES

REFERENCES

- [1] Mingyuan Fan, Shenqi Lai, Junshi Huang, Xiaoming Wei, Zhenhua Chai, Junfeng Luo, Xiaolin Wei. Rethinking BiSeNet For Real-time Semantic Segmentation. 2021
- [2] Changqian Yu, Jingbo Wang, Chao Peng, Changxin Gao, Gang Yu, and Nong Sang. Bisenet: Bilateral segmentation network for real-time semantic segmentation. In European conference on computer vision (ECCV) 2018.

TABLE IV
TRAIN ON GTA5 AUGEMENTED TRAIN SET AND TEST IT CITYSCAPES VAL SET WITHOUT ADAPTATION

GTA5 -> Cityscapes		
Accuracy (%)	mIoU (%)	Train time (avg per-epoch)
71	30	5.37 minutes

TABLE V
TRAIN ON GTA5 AUGEMENTED TRAIN SET AND TEST IT CITYSCAPES VAL SET WITH ADAPTATION

GTA5 -> Cityscapes		
Accuracy (%)	mIoU (%)	Train time (avg per-epoch)
74	33	4.30 minutes

TABLE VI

TRAIN ON GTA5 AUGEMENTED TRAIN SET AND TEST IT CITYSCAPES VAL
SET WITH LIGHTER ADAPTATION

GTA5 -> Cityscapes		
<i>Accuracy (%)</i>	<i>mIoU (%)</i>	<i>Train time (avg per-epoch)</i>
73	32.5	4.25 minutes

- [3] Yi-Hsuan Tsai, Wei-Chih Hung, Samuel Schulter, Kihyuk Sohn, Ming-Hsuan Yang, Manmohan Chandraker. Learning to Adapt Structured Output Space for Semantic Segmentation. In CVPR, 2018
- [4] S. R. Richter, V. Vineet, S. Roth, and V. Koltun. Playing for data: Ground truth from computer games. In ECCV, 2016.
- [5] M. Cordts, M. Omran, S. Ramos, T. Rehfeld, M. Enzweiler, R. Benenson, U. Franke, S. Roth, and B. Schiele. The cityscapes dataset for semantic urban scene understanding. In CVPR, 2016.
- [6] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, Hartwig Adam. MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications. 2017